

Vishal Gopinath

Bengaluru, Karnataka • +91 6362015896 • vishal.bglr@gmail.com •
linkedin.com/in/vishal-gopinath • vishu235.github.io/work-portfolio/

AI Engineer with 4+ years of experience building Generative AI, Deep Learning, and NLP systems in Python across finance and customer operations. Architected RAG knowledge platforms using OpenAI APIs and open-source LLMs with LangChain, LangGraph, and Weaviate, reducing hallucinated responses by 50%+. Designed multilingual Voice AI agents that automated ~70% of CSR calls across 500+ daily customer interactions. Deployed GPU-backed LLM and Machine Learning services on Kubernetes with Docker and LangSmith observability at enterprise scale.

WORK EXPERIENCE

MSI Services Private Limited AI Developer II

Bengaluru, India
04/2025 – Present

- Led enterprise-wide LLM initiatives across finance, operations, sales, and customer service, reducing manual effort by 40–60% across key workflows.
- Architected RAG-based knowledge systems using OpenAI APIs and Weaviate, and built customer-facing AI chatbots for product queries, order management, and delivery tracking — improving response turnaround by 70%+ and reducing hallucinations by 50%+.
- Designed multilingual Voice AI agents automating ~70% of CSR calls across English, French, and Spanish, and built an AI-driven call quality evaluation pipeline using STT and conversational analytics, eliminating manual QA for 10+ agents.
- Deployed GPU-backed LLM services on Kubernetes with LangSmith observability for production monitoring and cost optimization.

AI Developer

04/2023 – 03/2025

- Automated invoice and check deposit workflows using OCR + NLP extraction, improving efficiency by ~85% and reducing reconciliation errors.
- Built AI conversational systems (LangChain, LangGraph) handling 500+ daily customer queries with intelligent escalation, and developed a real-time RAG-powered assistant supporting CSRs during live calls.
- Architected self-learning product search and ranking pipeline processing 10,000+ orders/day, integrating AI services with enterprise backends via REST APIs.

Saigeware Technologies Private Limited AI Software Engineer

Bengaluru, India
07/2022 – 03/2023

- Developed DSP and speech enhancement pipelines using NumPy and SciPy, improving noise reduction accuracy from 63% to 81%, and implemented MFCC and PPG feature extraction to train ML models achieving ~90% classification accuracy.

EDUCATION

M. Tech. in Data Science & AI

PES University • GPA: 9.8

Bengaluru • 07/2024 – 08/2026

B. Tech. in Electronics and Communication Engineering

JSS Academy of Technical Education • GPA: 9.17

Bengaluru • 03/2018 – 06/2022

SKILLS

Core Expertise: Generative AI, RAG Systems, Conversational AI, Search & Ranking Models, NLP, Deep Learning, Multi-Agent Systems

LLMs & APIs: OpenAI APIs, Claude, Gemini, Open-source LLMs (LLaMA | Mistral | Falcon | Phi-3 | Gemma), Whisper

Data & Infrastructure: Python, Weaviate, Pinecone, Elasticsearch, MySQL, SQL, Docker, Kubernetes, REST APIs, LangChain, LangGraph, LangSmith, Weights & Biases

CERTIFICATIONS

Career Essentials in Generative AI by Microsoft and LinkedIn Microsoft	07/2024
Advance your skills in Natural Language Processing LinkedIn	08/2023
Advanced Python & DSA LinkedIn	03/2023
Prepare for Google Cloud Platform Associate Cloud Engineer LinkedIn	03/2023

PROJECTS

Meta-Cognitive Neuro-Symbolic Learning M.Tech Academic Project	04/2026 - 09/2026
- Developed an AI trust evaluation framework to detect when models produce correct predictions using biased patterns or shortcut reasoning, comparing multiple confidence estimation methods to identify overconfident or unreliable behavior across data shifts. - Proposed a meta-cognitive confidence layer that flags logically untrustworthy predictions, demonstrating that trustworthy AI requires both accurate outputs and reliable internal reasoning.	
AI Model to Analyze Sleep Quality M.Tech Academic Project	10/2025 - 03/2026
- Designed and implemented an end-to-end machine learning pipeline to predict sleep quality, performing feature engineering, model selection, and hyperparameter tuning across multiple classification algorithms to achieve 88% prediction accuracy. - Leveraged cloud GPU instances for scalable training and integrated Weights & Biases for experiment tracking, metric logging, and model performance monitoring across runs.	